

Programming for Data Science

Digital Assignment – 3

Feature Engineering & Model Fitting

Lab Slot: L33 & L34

Name: Tejas Rahul Rokade

Reg. No.: 20BDS0033

Questions:

1. Perform Label encoding on IRIS Dataset
2. Perform One-hot encoding on IRIS Dataset
3. Feature scaling or standardization
 - a. Normalization
 - b. Z-scale
4. Find the principal components of IRIS dataset
5. House rent prediction using linear regression
6. Medical diagnosis for disease spread pattern Using SVM

Q1 Perform Label Encoding on IRIS dataset.

Ans 1

Code:

```
library(datasets)
library(superml)
data = iris[c(1,2,3,5,46,48,49,90,98,101,102),]
label = LabelEncoder$new()
data$Species = label$fit_transform(data$Species)
data
```

Output:

```

C:\Users\Felipe Silva Mattar\AppData\Local\Temp\Rtmp3107\EO\downloaded_packages
> library(datasets)
> library(supernn)
Loading required package: R6
> data = iris[c(1,2,3,5,46,48,49,90,98,101,102),]
> label = LabelEncoder$new()
> data$species = label$fit_transform(data$species)
> data
  Sepal.Length Sepal.Width Petal.Length Petal.Width
1           5.1         3.5         1.4         0.2
2           4.9         3.0         1.4         0.2
3           4.7         3.2         1.3         0.2
5           5.0         3.6         1.4         0.2
46          4.8         3.0         1.4         0.3
48          4.6         3.2         1.4         0.2
49          5.3         3.7         1.5         0.2
90          5.5         2.5         4.0         1.3
98          6.2         2.9         4.3         1.3
101         6.3         3.3         6.0         2.5
102         5.8         2.7         5.1         1.9
  species
1       0
2       0
3       0
5       0
46      0
48      0
49      0
90      1
98      1
101     2
102     2
> |

```

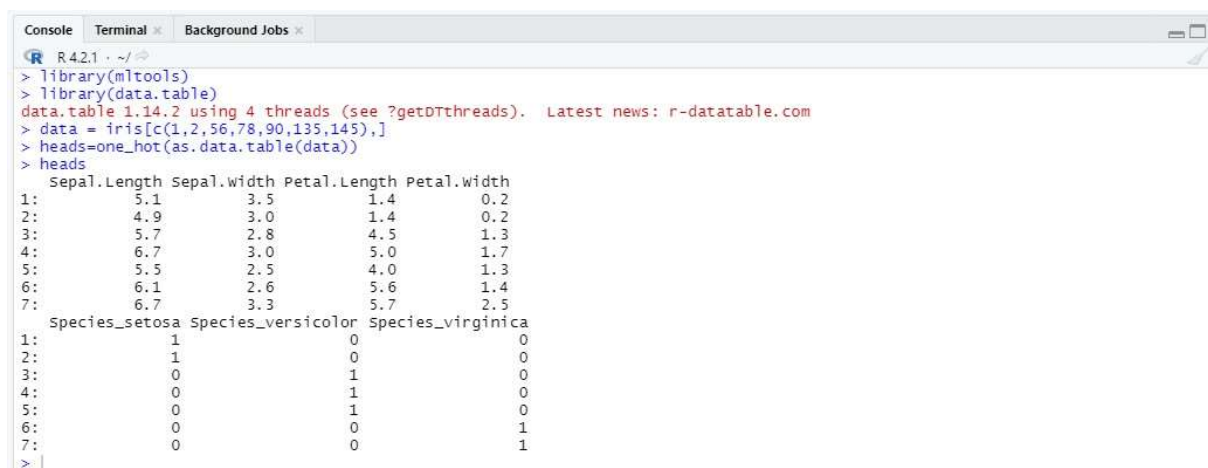
Q2. Perform One-hot encoding on IRIS Dataset

Ans 2.

Code:

```
library(mltools)
library(data.table)
data = iris[c(1,2,56,78,90,135,145),]
heads=one_hot(as.data.table(data))
heads
```

Output:



```
R 4.2.1 ~/>
> library(mltools)
> library(data.table)
data.table 1.14.2 using 4 threads (see ?getDTthreads). Latest news: r-datatable.com
> data = iris[c(1,2,56,78,90,135,145),]
> heads=one_hot(as.data.table(data))
> heads
   Sepal.Length Sepal.width Petal.Length Petal.width
1:         5.1         3.5         1.4         0.2
2:         4.9         3.0         1.4         0.2
3:         5.7         2.8         4.5         1.3
4:         6.7         3.0         5.0         1.7
5:         5.5         2.5         4.0         1.3
6:         6.1         2.6         5.6         1.4
7:         6.7         3.3         5.7         2.5
   Species_setosa Species_versicolor Species_virginica
1:              1              0              0
2:              1              0              0
3:              0              1              0
4:              0              1              0
5:              0              1              0
6:              0              0              1
7:              0              0              1
>
```

Q3. Feature scaling or standardization

a) Normalization

Code:

```
data = head(iris[-5])
data
y = sapply(data,function(data)((data-min(data))/(max(data)-min(data))))
y
```

Output:

```

R 4.2.1 ~ /
> data = head(iris[-5])
> data
  Sepal.Length Sepal.width Petal.Length Petal.width
1           5.1          3.5          1.4          0.2
2           4.9          3.0          1.4          0.2
3           4.7          3.2          1.3          0.2
4           4.6          3.1          1.5          0.2
5           5.0          3.6          1.4          0.2
6           5.4          3.9          1.7          0.4
> y = sapply(data,function(data)((data-min(data))/(max(data)-min(data))))
> y
  Sepal.Length Sepal.width Petal.Length Petal.width
[1,]      0.625 0.5555556      0.25      0
[2,]      0.375 0.0000000      0.25      0
[3,]      0.125 0.2222222      0.00      0
[4,]      0.000 0.1111111      0.50      0
[5,]      0.500 0.6666667      0.25      0
[6,]      1.000 1.0000000      1.00      1
>

```

b) Z-Scale

Code:

```
data = head(iris[-5])
```

```
data
```

```
y=sapply(data,function(data)((data-mean(data))/sd(data)))
```

```
y
```

Output:

```

R 4.2.1 ~ /
> data = head(iris[-5])
> data
  Sepal.Length Sepal.width Petal.Length Petal.width
1           5.1          3.5          1.4          0.2
2           4.9          3.0          1.4          0.2
3           4.7          3.2          1.3          0.2
4           4.6          3.1          1.5          0.2
5           5.0          3.6          1.4          0.2
6           5.4          3.9          1.7          0.4
> y=sapply(data,function(data)((data-mean(data))/sd(data)))
> y
  Sepal.Length Sepal.width Petal.Length Petal.width
[1,]  0.5206576  0.3401105 -0.3627381 -0.4082483
[2,] -0.1735525 -1.1175060 -0.3627381 -0.4082483
[3,] -0.8677627 -0.5344594 -1.0882144 -0.4082483
[4,] -1.2148677 -0.8259827  0.3627381 -0.4082483
[5,]  0.1735525  0.6316338 -0.3627381 -0.4082483
[6,]  1.5619728  1.5062037  1.8136906  2.0412415
>

```

Q4. Find the principal components of IRIS dataset.

Ans 4

Code:

```
library(dplyr)
library(datasets)
library(CatEncoders)
library(superml)
data = iris[-5]
my_pca = prcomp(data, scale. = TRUE, center = TRUE, retx = T)
names(my_pca)
summary(my_pca)
my_pca$rotation
my_pca$sdev
```

Output:

```

> library(datasets)
> library(CatEncoders)

Attaching package: 'CatEncoders'

The following object is masked from 'package:base':

    transform

> library(superml)
> data = iris[-5]
> my_pca = prcomp(data, scale. = TRUE, center = TRUE, retx = T)
> names(my_pca)
[1] "sdev"      "rotation" "center"   "scale"
[5] "x"
> summary(my_pca)
Importance of components:
              PC1      PC2      PC3      PC4
Standard deviation   1.7084 0.9560 0.38309 0.14393
Proportion of Variance 0.7296 0.2285 0.03669 0.00518
Cumulative Proportion 0.7296 0.9581 0.99482 1.00000
> my_pca$rotation
              PC1      PC2      PC3
Sepal.Length  0.5210659 -0.37741762 0.7195664
Sepal.Width   -0.2693474 -0.92329566 -0.2443818
Petal.Length  0.5804131 -0.02449161 -0.1421264
Petal.Width   0.5648565 -0.06694199 -0.6342727
              PC4
Sepal.Length  0.2612863
Sepal.Width   -0.1235096
Petal.Length  -0.8014492
Petal.Width   0.5235971
> my_pca$sdev
[1] 1.7083611 0.9560494 0.3830886 0.1439265
> |

```

Q5 House rent prediction using linear regression**Ans 5****Code:**

```
library(mlbench)

library(caTools)

data(BostonHousing)

x=BostonHousing

str(x)

sum(is.na(x))

split=sample.split(x,SplitRatio = 0.8)

train=subset(x,split==TRUE)

test=subset(x,split==FALSE)

model=lm(medv ~ crim + rm + tax + lstat, data=train)

summary(model)

test$predicted.medv=predict(model,test)

print(test$medv)

print(test$predicted.medv)

error=test$medv-test$predicted.medv

rmse=sqrt(mean(error)^2)

cat("RMSE",rmse)
```

Output:

```

Console Terminal x Background Jobs x
R 4.2.1 ~ /
> library(mlbench)
> library(caTools)
> data(BostonHousing)
> x=BostonHousing
> str(x)
'data.frame':  506 obs. of  14 variables:
 $ crim   : num  0.00632 0.02731 0.02729 0.03237 0.06905 ...
 $ zn     : num  18 0 0 0 0 12.5 12.5 12.5 12.5 ...
 $ indus  : num  2.31 7.07 7.07 2.18 2.18 2.18 7.87 7.87 7.87 ...
 $ chas   : Factor w/ 2 levels "0","1": 1 1 1 1 1 1 1 1 1 ...
 $ nox    : num  0.538 0.469 0.469 0.458 0.458 0.458 0.524 0.524 0.524 ...
 $ rm     : num  6.58 6.42 7.18 7 7.15 ...
 $ age    : num  65.2 78.9 61.1 45.8 54.2 58.7 66.6 96.1 100 85.9 ...
 $ dis    : num  4.09 4.97 4.97 6.06 6.06 ...
 $ rad    : num  1 2 2 3 3 3 5 5 5 ...
 $ tax    : num  296 242 242 222 222 222 311 311 311 311 ...
 $ ptratio: num  15.3 17.8 17.8 18.7 18.7 18.7 15.2 15.2 15.2 15.2 ...
 $ b      : num  397 397 393 395 397 ...
 $ lstat  : num  4.98 9.14 4.03 2.94 5.33 ...
 $ medv   : num  24 21.6 34.7 33.4 36.2 28.7 22.9 27.1 16.5 18.9 ...
> sum(is.na(x))
[1] 0
> split=sample.split(x,SplitRatio = 0.8)
> train=subset(x,split==TRUE)
> test=subset(x,split==FALSE)
> model=lm(medv ~ crim + rm + tax + lstat, data=train)
> summary(model)

Call:
lm(formula = medv ~ crim + rm + tax + lstat, data = train)

Residuals:
    Min       1Q   Median       3Q      Max
-16.433  -3.386  -1.103   1.873  30.878

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept) -1.735635   3.380444  -0.513   0.6079
crim         -0.063053   0.040171  -1.570   0.1173
rm           5.304867   0.465932  11.385 <2e-16 ***
tax          -0.005176   0.002116  -2.446   0.0149 *
lstat        -0.537327   0.054418  -9.874 <2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 5.344 on 393 degrees of freedom
Multiple R-squared:  0.6627,    Adjusted R-squared:  0.6593
F-statistic: 193 on 4 and 393 DF,  p-value: < 2.2e-16

```



```

> test$predicted.medv=predict(model,test)
> print(test$medv)
[1] 22.9 27.1 20.4 13.6 19.6 14.8 13.5 18.9 26.6 14.4 19.4 35.4 22.2
[14] 25.0 20.9 20.0 20.8 22.9 22.6 22.0 38.7 20.1 19.5 22.8 20.4 19.3
[27] 21.4 23.0 18.4 17.8 15.6 14.6 19.4 27.0 50.0 23.8 22.6 29.4 36.2
[40] 29.8 34.9 50.0 42.3 48.5 20.0 23.3 28.7 30.1 24.3 31.7 31.5 17.6
[53] 18.5 24.8 36.0 30.1 22.8 24.4 35.2 35.1 20.1 23.2 23.9 24.8 22.0
[66] 28.2 23.8 16.2 23.1 19.3 22.6 21.1 16.5 23.9 26.6 17.8 21.7 16.8
[79] 50.0 50.0 13.3 8.8 7.2 23.2 5.0 6.3 5.0 17.9 16.3 8.4 10.2
[92] 10.9 14.3 10.5 17.1 12.6 14.9 14.1 17.7 19.1 20.1 13.3 25.0 21.8
[105] 7.0 19.7 18.3 23.9
> print(test$predicted.medv)
[1] 21.8630337 19.0974292 23.7560410 14.8499860 20.8339933
[6] 19.4086746 17.9829639 23.0887420 30.3636711 9.1283511
[11] 18.0577331 32.9642611 27.4195158 27.5533866 22.9666963
[16] 23.0754802 23.2524709 25.4875355 26.1714461 26.4360311
[21] 37.3710693 22.3579572 18.4575267 26.1850322 18.9116820
[26] 19.1084967 21.0777939 23.7929042 18.7906883 18.6794211
[31] 16.9519625 6.2937111 17.8441279 26.2983703 34.8848482
[36] 20.7185875 22.6284075 28.5898536 24.7764697 28.5249899
[41] 31.4184238 37.1236486 35.1607885 36.7146923 12.7455483
[46] 20.8087596 26.8596706 27.6605821 22.1101079 33.1449320
[51] 33.1347976 19.4969157 16.3590367 28.7666045 31.5709485
[56] 29.4449123 20.7299444 27.6326346 34.3603074 30.6638600
[61] 23.1174603 25.8468403 24.6560230 29.5968351 26.4137878
[66] 29.3992054 26.5183318 20.7599627 26.8998766 21.8064397
[71] 25.6863046 24.8239975 26.1245845 28.0111236 30.1742837
[76] 17.7478043 21.3688895 17.4687713 30.0334961 22.0895289
[81] 18.8263690 0.2642467 5.1960320 16.5030833 4.8897290
[86] 9.1324542 8.3344773 -0.2857662 9.5675492 18.0541235
[91] 16.5789517 17.5414157 19.8921063 12.4103517 17.6985446
[96] 18.5186037 19.8552222 19.3901977 20.5364597 15.5301342
[101] 16.3686602 14.1532354 28.1469637 19.6073184 10.4136859
[106] 13.4565389 19.3838235 30.8237875
> error=test$medv-test$predicted.medv
> rmse=sqrt(mean(error)^2)
> cat("RMSE",rmse)
RMSE 0.322775
> |

```

Q6 Medical diagnosis for disease spread pattern Using SVM

Ans 6

Code:

```
library(superml)
```

```
library(caTools)
```

```
library(e1071)
```

```
library(tidyverse)
```

```
x=read.csv("Cancer_Data.CSV")
```

```
names(x)
```

```
x=x[-c(1,33)]
```

```
sum(is.na(x))
```

```
colSums(is.na(x))
```

```
label=LabelEncoder$new()
```

```
x$diagnosis=label$fit_transform(x$diagnosis)
```

```

head(x)

split=sample.split(x$diagnosis,SplitRatio = 0.8)

train=subset(x,split==TRUE)

test=subset(x,split==FALSE)

train[-1]=scale(train[-1])

test[-1]=scale(test[-1])

names(train)

classifier=svm(formula=diagnosis~.,
               data = train,
               type = 'C-Classification',
               kernel = 'linear')

Diag_pred=predict(classifier, newdata = test[-1])

cm = table(test[,1], Daig_pred)

print(cm)

```

Output:

```

> library(supernl)
> library(caTools)
> library(e1071)
> x=read.csv("Cancer_Data.csv")
> names(x)
 [1] "id" "diagnosis" "radius_mean" "texture_mean"
 [5] "perimeter_mean" "area_mean" "smoothness_mean" "compactness_mean"
 [9] "concavity_mean" "concave.points_mean" "symmetry_mean" "fractal_dimension_mean"
[13] "radius_se" "texture_se" "perimeter_se" "area_se"
[17] "smoothness_se" "compactness_se" "concavity_se" "concave.points_se"
[21] "symmetry_se" "fractal_dimension_se" "radius_worst" "texture_worst"
[25] "perimeter_worst" "area_worst" "smoothness_worst" "compactness_worst"
[29] "concavity_worst" "concave.points_worst" "symmetry_worst" "fractal_dimension_worst"
[33] "x"
> x=x[-c(1,33)]
> sum(is.na(x))
[1] 0

```

```

> colSums(is.na(x))
      diagnosis      radius_mean      texture_mean      perimeter_mean
      0              0              0              0
      area_mean      smoothness_mean      compactness_mean      concavity_mean
      0              0              0              0
      concave.points_mean      symmetry_mean      fractal_dimension_mean      radius_se
      0              0              0              0
      texture_se      perimeter_se      area_se      smoothness_se
      0              0              0              0
      compactness_se      concavity_se      concave.points_se      symmetry_se
      0              0              0              0
      fractal_dimension_se      radius_worst      texture_worst      perimeter_worst
      0              0              0              0
      area_worst      smoothness_worst      compactness_worst      concavity_worst
      0              0              0              0
      concave.points_worst      symmetry_worst      fractal_dimension_worst
      0              0              0
> label=LabelEncoder$new()
> x$diagnosis=label$fit_transform(x$diagnosis)

> head(x)
diagnosis radius_mean texture_mean perimeter_mean area_mean smoothness_mean compactness_mean concavity_mean
1      0      17.99      10.38      122.80      1001.0      0.11840      0.27760      0.3001
2      0      20.57      17.77      132.90      1326.0      0.08474      0.07864      0.0869
3      0      19.69      21.25      130.00      1203.0      0.10960      0.15990      0.1974
4      0      11.42      20.38      77.58      386.1      0.14250      0.28390      0.2414
5      0      20.29      14.34      135.10      1297.0      0.10030      0.13280      0.1980
6      0      12.45      15.70      82.57      477.1      0.12780      0.17000      0.1578
concave.points_mean symmetry_mean fractal_dimension_mean radius_se texture_se perimeter_se area_se smoothness_se
1      0.14710      0.2419      0.07871      1.0950      0.9053      8.589      153.40      0.006399
2      0.07017      0.1812      0.05667      0.5435      0.7339      3.398      74.08      0.005225
3      0.12790      0.2069      0.05999      0.7456      0.7869      4.585      94.03      0.006150
4      0.10520      0.2597      0.09744      0.4956      1.1560      3.445      27.23      0.009110
5      0.10430      0.1809      0.05883      0.7572      0.7813      5.438      94.44      0.011490
6      0.08089      0.2087      0.07613      0.3345      0.8902      2.217      27.19      0.007510
compactness_se concavity_se concave.points_se symmetry_se fractal_dimension_se radius_worst texture_worst
1      0.04904      0.05373      0.01587      0.03003      0.006193      25.38      17.33
2      0.01308      0.01860      0.01340      0.01389      0.003532      24.99      23.41
3      0.04006      0.03832      0.02058      0.02250      0.004571      23.57      25.53
4      0.07458      0.05661      0.01867      0.05963      0.009208      14.91      26.50
5      0.02461      0.05688      0.01885      0.01756      0.005115      22.54      16.67
6      0.03345      0.03672      0.01137      0.02165      0.005082      15.47      23.75
perimeter_worst area_worst smoothness_worst compactness_worst concavity_worst concave.points_worst symmetry_worst
1      184.60      2019.0      0.1622      0.6656      0.7119      0.2654      0.4601
2      158.80      1956.0      0.1238      0.1866      0.2416      0.1860      0.2750
3      152.50      1709.0      0.1444      0.4245      0.4504      0.2430      0.3613
4      98.87      567.7      0.2098      0.8663      0.6869      0.2575      0.6638
5      152.20      1575.0      0.1374      0.2050      0.4000      0.1625      0.2364
6      103.40      741.6      0.1791      0.5249      0.5355      0.1741      0.3985

fractal_dimension_worst
1      0.11890
2      0.08902
3      0.08758
4      0.17300
5      0.07678
6      0.12440

```

```

> split=sample.split(x$diagnosis,SplitRatio =0.8)
> train=subset(x,split==TRUE)
> test=subset(x,split==FALSE)
> train[-1]=scale(train[-1])
> test[-1]=scale(test[-1])
> names(train)
[1] "diagnosis"      "radius_mean"    "texture_mean"    "perimeter_mean"
[5] "area_mean"      "smoothness_mean" "compactness_mean" "concavity_mean"
[9] "concave.points_mean" "symmetry_mean"  "fractal_dimension_mean" "radius_se"
[13] "texture_se"      "perimeter_se"   "area_se"         "smoothness_se"
[17] "compactness_se"  "concavity_se"   "concave.points_se" "symmetry_se"
[21] "fractal_dimension_se" "radius_worst"  "texture_worst"    "perimeter_worst"
[25] "area_worst"      "smoothness_worst" "compactness_worst" "concavity_worst"
[29] "concave.points_worst" "symmetry_worst" "fractal_dimension_worst"
> classifier = svm(formula = diagnosis ~ .,
+                   data = train,
+                   type = 'C-classification',
+                   kernel = 'linear')
> Diag_pred = predict(classifier, newdata = test[-1])

> cm = table(test[,1], Diag_pred)
> print(cm)
  Diag_pred
    0    1
0  41   1
1   1  70

```