

Programming For Data Science Lab Assignment 4

L33-L34

Tejas Rokade

20BDS0033

Q1. Read gender_classification dataset from R and perform following model fitting techniques (Any other datasets also applicable)

- a. Logistic Regression**
- b. Decision Tree**
- c. Naïve Bayes**
- d. SVM**
- e. Random forest**

Q2. Compare each of the above models using the following parameters

- a. Accuracy**
- b. Precision**
- c. Recall**
- d. Sensitivity**
- e. Specificity**

Ans 1 & 2.

Code:

```
#FUNCTION TO FIND accuracy, precision, recall, Sensitivity, Specificity
```

```
classification_report_for_model <- function(conf_matr){
```

```
print(paste0("Accuracy: ", (
```

```
(conf_matr[1,1]+conf_matr[2,2])/(conf_matr[1,1]+conf_matr[1,2]+conf_matr[2,1]+conf_m  
atr[2,2]) )))
```

```
print("")
```

```

print(paste0("Precision: ", (conf_matr[2,2]/(conf_matr[2,2]+conf_matr[1,2])))
print("")
print(paste0("Recall: ", (conf_matr[2,2]/(conf_matr[2,2]+conf_matr[2,1])) ))
print("")
print(paste0("Sensitivity: ", (conf_matr[2,2]/(conf_matr[2,2]+conf_matr[2,1])) ))
print("")
print(paste0("Specificity: ", (conf_matr[1,1]/(conf_matr[1,1]+conf_matr[1,2])))

}

Gender_class=read.csv("C:\\Users\\Parshva Maniar\\Downloads\\archive\\Transformed
Data Set - Sheet1.csv")

library(superml)

label=LabelEncoder$new()

Gender_class$Favorite.Color=label$fit_transform(Gender_class$Favorite.Color)

label=LabelEncoder$new()

Gender_class$Favorite.Music.Genre=label$fit_transform(Gender_class$Favorite.Music.
Genre)

label=LabelEncoder$new()

Gender_class$Favorite.Beverage=label$fit_transform(Gender_class$Favorite.Beverage)

label=LabelEncoder$new()

Gender_class$Favorite.Soft.Drink =label$fit_transform(Gender_class$Favorite.Soft.Drink
)

label=LabelEncoder$new()

Gender_class$Gender=label$fit_transform(Gender_class$Gender)

Gender_class$Gender = factor(Gender_class$Gender, levels = c(0, 1))

head(Gender_class)

library(caTools)

#splitting data into train and test

```

```

split=sample.split(Gender_class$Gender,SplitRatio =0.7)
train=subset(Gender_class,split==TRUE)
test=subset(Gender_class,split==FALSE)
library(e1071)
train[-5]=scale(train[-5])
test[-5]=scale(test[-5])
#LOGISTIC
model_log_reg =glm(Gender ~ ., data = train,family = "binomial")
model_log_reg
#Prediction
pred_lr = predict(model_log_reg, newdata = test[-5],type="response")
pred_lr <- ifelse(pred_lr>0.5,1,0)
pred_lr
# Making the Confusion Matrix
conf_matr = table(test[,5], pred_lr)
print(conf_matr)
classification_report_for_model(conf_matr)
#Desicion tree
library(party)
Desc_tree = ctree(Gender~., data = train)
pred_desc_tree = predict(Desc_tree, newdata = test[-5])
pred_desc_tree
# Making the Confusion Matrix
conf_matr = table(test[,5], pred_desc_tree)
print(conf_matr)
classification_report_for_model(conf_matr)
#Naive Bayes
naive_bayes =naiveBayes(Gender ~ ., data = train)

```

```

naive_bayes
#Prediction
pred_nb = predict(naive_bayes, newdata = test[-5])
# Making the Confusion Matrix
conf_matr = table(test[,5], pred_nb)
print(conf_matr)
classification_report_for_model(conf_matr)
#SVM
svm_model = svm(formula = Gender ~ ., data = train, type = 'C-classification', kernel
                = 'linear')
svm_model
#Prediction
pred_svm = predict(svm_model, newdata = test[-5])
#Confusion Matrix
conf_matr = table(test[,5], pred_svm)
print(conf_matr)
classification_report_for_model(conf_matr)
#Random Forest
library(randomForest)
randomforest_model = randomForest(x =train[-5], y =train$Gender, ntree = 500)
randomforest_model
#predict
pred_rft = predict(randomforest_model, newdata = test[-5])
#Confusion Matrix
conf_matr = table(test[,5], pred_rft)
conf_matr
classification_report_for_model(conf_matr)

```

Output:

```

R 4.2.1 ~ /
> #FUNCTION TO FIND accuracy, precision, recall, sensitivity, specificity
> classification_report_for_model <- function(conf_matr){
+   print(paste0("Accuracy: ", (
+     (conf_matr[1,1]+conf_matr[2,2])/(conf_matr[1,1]+conf_matr[1,2]+conf_matr[2,1]+conf_matr[2,2]) )))
+   print("")
+   print(paste0("Precision: ", (conf_matr[2,2]/(conf_matr[2,2]+conf_matr[1,2]))))
+   print("")
+   print(paste0("Recall: ", (conf_matr[2,2]/(conf_matr[2,2]+conf_matr[2,1])) ))
+   print("")
+   print(paste0("Sensitivity: ", (conf_matr[2,2]/(conf_matr[2,2]+conf_matr[2,1])) ))
+   print("")
+   print(paste0("Specificity: ", (conf_matr[1,1]/(conf_matr[1,1]+conf_matr[1,2]))))
+ }
> Gender_class=read.csv("C:\\Users\\Parshva Maniar\\Downloads\\archive\\Transformed Data Set - Sheet1.csv")
> library(supernl)
> label=LabelEncoder$new()
> Gender_class$Favorite.color=label$fit_transform(Gender_class$Favorite.color)
> label=LabelEncoder$new()
> Gender_class$Favorite.Music.Genre=label$fit_transform(Gender_class$Favorite.Music.Genre)
> label=LabelEncoder$new()
> Gender_class$Favorite.Beverage=label$fit_transform(Gender_class$Favorite.Beverage)
> label=LabelEncoder$new()
> Gender_class$Favorite.Soft.Drink =label$fit_transform(Gender_class$Favorite.Soft.Drink
+ )
> label=LabelEncoder$new()
> Gender_class$Gender=label$fit_transform(Gender_class$Gender)
> Gender_class$Gender = factor(Gender_class$Gender, levels = c(0, 1))
> head(Gender_class)
  Favorite.Color Favorite.Music.Genre Favorite.Beverage Favorite.Soft.Drink Gender
1             0                  0              0              0             0
2             1                  1              0              1             0
3             2                  0              1              1             0
4             2                  2              2              2             0
5             0                  0              0              1             0
6             2                  3              3              2             0

```

```

R 4.2.1 ~ /
> library(caTools)
> #splitting data into train and test
> split=sample.split(Gender_class$Gender,splitRatio =0.7)
> train=subset(Gender_class,split==TRUE)
> test=subset(Gender_class,split==FALSE)
> library(e1071)
> train[-5]=scale(train[-5])
> test[-5]=scale(test[-5])
> #LOGISTIC
> model_log_reg =glm(Gender ~ ., data = train,family = "binomial")
> model_log_reg

Call: glm(formula = Gender ~ ., family = "binomial", data = train)

Coefficients:
(Intercept)      Favorite.Color  Favorite.Music.Genre  Favorite.Beverage  Favorite.Soft.Drink
    0.01176         -0.24580         -0.10460          0.06401          0.62048

Degrees of Freedom: 45 Total (i.e. Null);  41 Residual
Null Deviance:      63.77
Residual Deviance: 59.83    AIC: 69.83
> #Prediction
> pred_lr = predict(model_log_reg, newdata = test[-5],type="response")
> pred_lr <- ifelse(pred_lr>0.5,1,0)
> pred_lr
 4  9 12 14 15 21 27 28 29 33 38 39 41 42 43 51 53 55 64 65
0  0  0  0  1  1  1  0  0  1  1  0  0  0  0  0  1  1  1  0
> # Making the Confusion Matrix
> conf_matr = table(test[,5], pred_lr)
> print(conf_matr)
  pred_lr
    0  1
0  6  4
1  5  5
> classification_report_for_model(conf_matr)
[1] "Accuracy: 0.55"
[1] ""
[1] "Precision: 0.555555555555556"
[1] ""
[1] "Recall: 0.5"
[1] ""
[1] "Sensitivity: 0.5"
[1] ""
[1] "Specificity: 0.6"
> #Desicion tree
> library(party)
> Desc_tree = ctree(Gender~., data = train)
> pred_desc_tree = predict(Desc_tree, newdata = test[-5])
> pred_desc_tree
[1] 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
Levels: 0 1

```

```

R 4.2.1 ~ /
Console Terminal Background Jobs x
> conf_matr = table(test[,5], pred_desc_tree)
> print(conf_matr)
  pred_desc_tree
0 1
0 10 0
1 10 0
> classification_report_for_model(conf_matr)
[1] "Accuracy: 0.5"
[1] ""
[1] "Precision: NaN"
[1] ""
[1] "Recall: 0"
[1] ""
[1] "Sensitivity: 0"
[1] ""
[1] "Specificity: 1"
> #Naive Bayes
> naive_bayes = naiveBayes(Gender ~ ., data = train)
> naive_bayes

Naive Bayes Classifier for Discrete Predictors

Call:
naiveBayes.default(x = x, y = y, laplace = laplace)

A-priori probabilities:
y
 0 1
0.5 0.5

Conditional probabilities:
Favorite.Color
y      [,1]      [,2]
0 0.06922995 1.007715
1 -0.06922995 1.009923

Favorite.Music.Genre
y      [,1]      [,2]
0 0.02049915 0.950848
1 -0.02049915 1.067925

Favorite.Beverage
y      [,1]      [,2]
0 -0.07356222 1.0306934
1 0.07356222 0.9858047

Favorite.Soft.Drink
y      [,1]      [,2]
0 -0.2530263 0.7778618
1 0.2530263 1.1430314

> #Prediction
> pred_nb = predict(naive_bayes, newdata = test[-5])
> # Making the Confusion Matrix
> conf_matr = table(test[,5], pred_nb)
> print(conf_matr)
  pred_nb
0 1
0 7 3
1 5 5
> classification_report_for_model(conf_matr)
[1] "Accuracy: 0.6"
[1] ""
[1] "Precision: 0.625"
[1] ""
[1] "Recall: 0.5"
[1] ""
[1] "Sensitivity: 0.5"
[1] ""
[1] "Specificity: 0.7"
> #SVM
> svm_model = svm(formula = Gender ~ ., data = train, type = 'c-classification', kernel
+                 = 'linear')
> svm_model

Call:
svm(formula = Gender ~ ., data = train, type = "C-classification", kernel = "linear")

Parameters:
  SVM-Type:  C-classification
 SVM-Kernel: linear
      cost:  1

Number of Support Vectors: 42

> #Prediction
> pred_svm = predict(svm_model, newdata = test[-5])
> #Confusion Matrix
> conf_matr = table(test[,5], pred_svm)
> print(conf_matr)
  pred_svm
0 1
0 7 3
1 6 4

```

```

> classification_report_for_model(conf_matr)
[1] "Accuracy: 0.55"
[1] ""
[1] "Precision: 0.571428571428571"
[1] ""
[1] "Recall: 0.4"
[1] ""
[1] "Sensitivity: 0.4"
[1] ""
[1] "Specificity: 0.7"
> #Random Forest
> library(randomForest)
randomForest 4.7-1.1
Type rfNews() to see new features/changes/bug fixes.
> randomforest_model = randomForest(x =train[-5], y =train$gender, ntree = 500)
> randomforest_model

Call:
randomForest(x = train[-5], y = train$gender, ntree = 500)
Type of random forest: classification
Number of trees: 500
No. of variables tried at each split: 2

OOB estimate of error rate: 41.3%
Confusion matrix:
  0  1 class.error
0 16  7  0.3043478
1 12 11  0.5217391
> #predict
> pred_rft = predict(randomforest_model, newdata = test[-5])
> #Confusion Matrix
> conf_matr = table(test[,5], pred_rft)
> conf_matr
  pred_rft
  0  1
0  6  4
1  6  4
> classification_report_for_model(conf_matr)
[1] "Accuracy: 0.5"
[1] ""
[1] "Precision: 0.5"
[1] ""
[1] "Recall: 0.4"
[1] ""
[1] "Sensitivity: 0.4"
[1] ""
[1] "Specificity: 0.6"
>

```

Q3. Perform K Means Clustering on IRIS dataset

Ans 3.

CODE:

Removing initial label of

```
x <- iris[, -5]
```

```
head(x)
```

Fitting the clustering Model to the dataset

```
set.seed(240) # Setting seed
```

```
kmeans_clust <- kmeans(x, centers = 3, nstart = 20)
```

```
kmeans_clust
```

```
kmeans_clust$cluster
```

Confusion Matrix

```
cm <- table(iris$Species, kmeans_clust$cluster)
```


```
cm
```

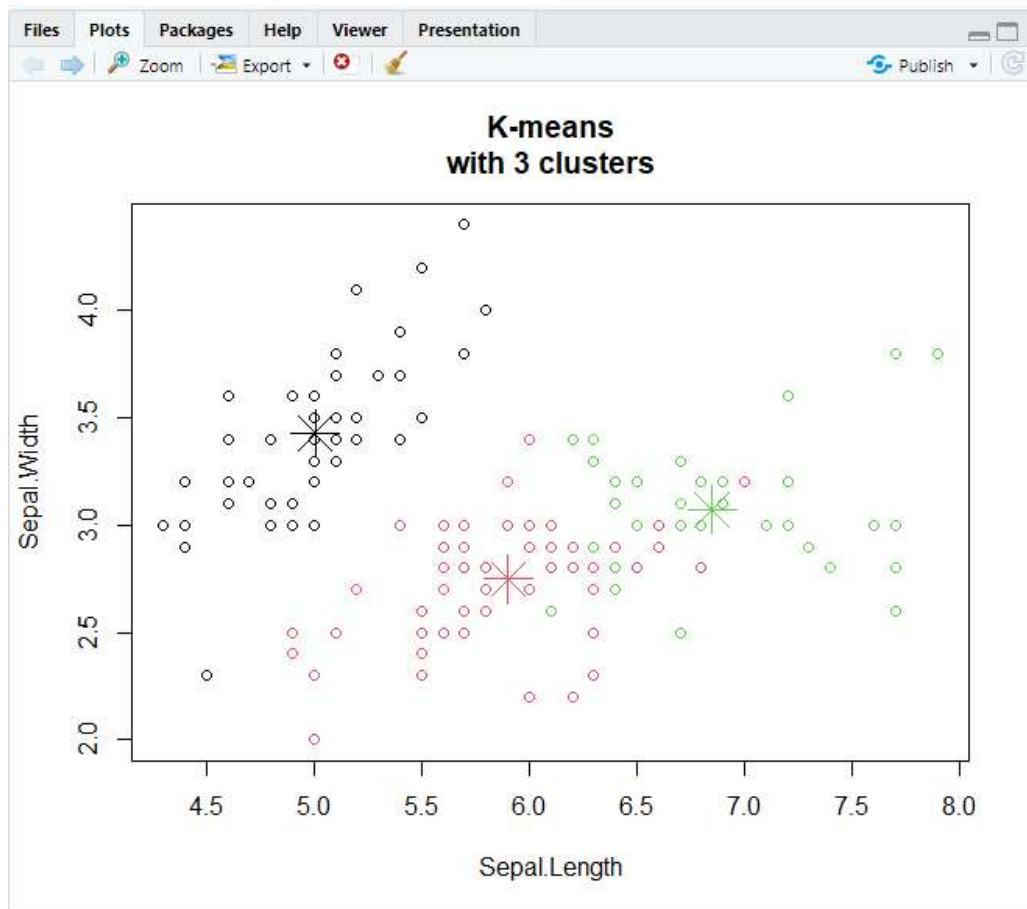
Visualization

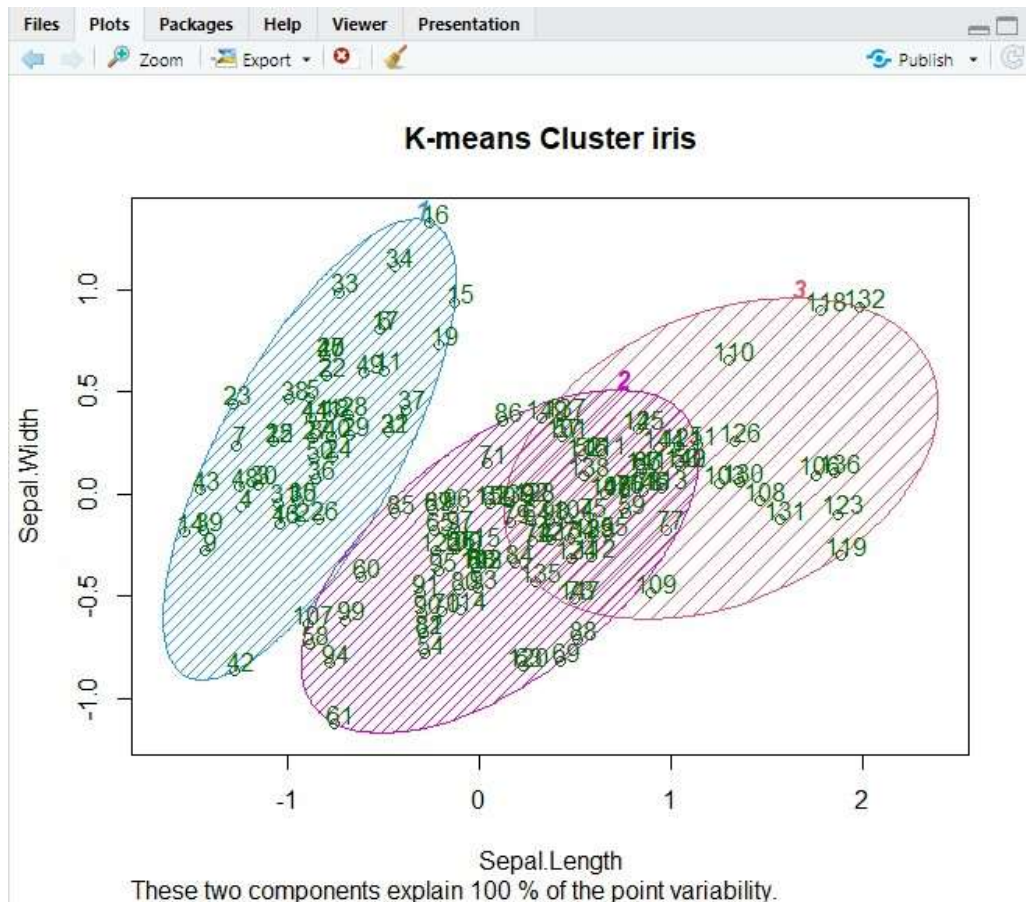
```
plot(x[c("Sepal.Length", "Sepal.Width")], col = kmeans_clust$cluster, main = "K-means
with 3 clusters")
# Plotting clusters with centres
kmeans_clust$centers
kmeans_clust$centers[, c("Sepal.Length", "Sepal.Width")]
points(kmeans_clust$centers[, c("Sepal.Length", "Sepal.Width")], col = 1:3, pch = 8,
       cex = 3)
# Visualizing clusters
y <- kmeans_clust$cluster
library(cluster)
clusplot(x[, c("Sepal.Length", "Sepal.Width")], y, lines = 0, shade = TRUE, color = TRUE,
        labels = 2, plotchar = FALSE, span = TRUE, main = paste("K-means Cluster iris"),
        xlab
        = 'Sepal.Length', ylab = 'Sepal.Width')
```

Output:

[illegible]







Q4. Perform Hierarchical clustering in mtcars dataset

Ans 4

Code:

```
mtcars_temp <- dist(mtcars, method = 'euclidean')
mtcars_temp
set.seed(240)
hierarchial_cluster_model <- hclust(mtcars_temp, method = "average")
hierarchial_cluster_model
# Plotting dendrogram
plot(hierarchial_cluster_model)
# Cutting tree by height
```

```
abline(h = 110, col = "green")
```

```
# Cutting by no. of clusters
```

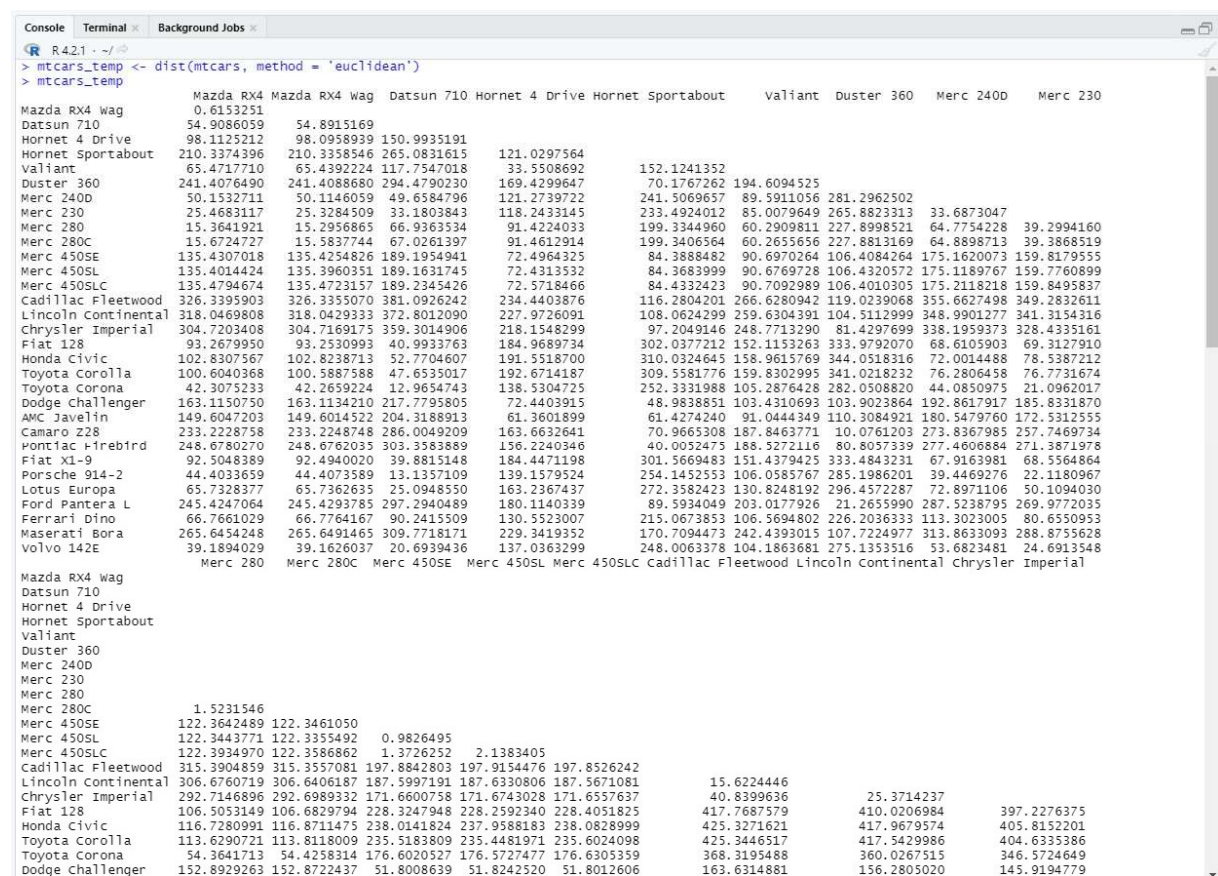
```
final_fit <- cutree(hierarchical_cluster_model, k = 3)
```

```
final_fit
```

```
table(final_fit)
```

```
rect.hclust(hierarchical_cluster_model, k = 3, border = "yellow")
```

Output:



```
R 4.2.1 ~ />  
> mtcars_temp <- dist(mtcars, method = 'euclidean')  
> mtcars_temp
```

| | Mazda RX4 | Mazda RX4 Wag | Datsun 710 | Hornet 4 Drive | Hornet Sportabout | valiant | Duster 360 | Merc 240D | Merc 230 |
|---------------------|-------------|---------------|-------------|----------------|-------------------|-------------|-------------|-------------|-------------|
| Mazda RX4 Wag | 0.6153251 | | | | | | | | |
| Datsun 710 | 54.9086059 | 54.8915169 | | | | | | | |
| Hornet 4 Drive | 98.1125212 | 98.0958939 | 150.9935191 | | | | | | |
| Hornet Sportabout | 210.3374396 | 210.3358546 | 265.0831615 | 121.0297564 | | | | | |
| valiant | 65.4717710 | 65.4392224 | 117.7547018 | 33.5508692 | 152.1241352 | | | | |
| Duster 360 | 241.4076490 | 241.4088680 | 294.4790230 | 169.4299647 | 70.1767262 | 194.6094525 | | | |
| Merc 240D | 50.1532711 | 50.1146059 | 49.6584796 | 121.2739722 | 241.5069657 | 89.5911056 | 281.2962502 | | |
| Merc 230 | 25.4683117 | 25.3284509 | 33.1803843 | 118.2433145 | 233.4924012 | 85.0079649 | 265.8823313 | 33.6873047 | |
| Merc 280 | 15.3641921 | 15.2956865 | 66.9363534 | 91.4224033 | 199.3344960 | 60.2909811 | 227.8998521 | 64.7754228 | 39.2994160 |
| Merc 280C | 15.6724727 | 15.5837744 | 67.0261397 | 91.4612914 | 199.3406564 | 60.2655656 | 227.8813169 | 64.8898713 | 39.3868519 |
| Merc 450SE | 135.4307018 | 135.4254826 | 189.1954941 | 72.4964325 | 84.3888482 | 90.6970264 | 106.4084264 | 175.1620073 | 159.8179555 |
| Merc 450SL | 135.4014424 | 135.3960351 | 189.1631745 | 72.4313532 | 84.3683999 | 90.6769728 | 106.4320572 | 175.1189767 | 159.7760899 |
| Merc 450SLC | 135.4794674 | 135.4723157 | 189.2345426 | 72.5718466 | 84.4332423 | 90.7092989 | 106.4010305 | 175.2118218 | 159.8495837 |
| Cadillac Fleetwood | 326.3395903 | 326.3355070 | 381.0926242 | 234.4403876 | 116.2804201 | 266.6280942 | 119.0239068 | 355.6627498 | 349.2832611 |
| Lincoln Continental | 318.0469808 | 318.0429333 | 372.8012090 | 227.9726091 | 108.0624299 | 259.6304391 | 104.5112999 | 348.9901277 | 341.3154316 |
| Chrysler Imperial | 304.7203408 | 304.7169175 | 359.3014906 | 218.1548299 | 97.2049146 | 248.7713290 | 81.4297699 | 338.1959373 | 328.4335161 |
| Fiat 128 | 93.2679950 | 93.2530993 | 40.9933763 | 184.9689734 | 302.0377212 | 152.1153263 | 333.9792070 | 68.6105903 | 69.3127910 |
| Honda Civic | 102.8307567 | 102.8238713 | 52.7704607 | 191.5518700 | 310.0324645 | 158.9615769 | 344.0518316 | 72.0014488 | 78.5387212 |
| Toyota Corolla | 100.6040368 | 100.5887588 | 47.6535017 | 192.6714187 | 309.5581776 | 159.8302995 | 341.0218232 | 76.2806458 | 76.7731674 |
| Toyota Corona | 42.3075233 | 42.2659224 | 12.9654743 | 138.5304725 | 252.3331988 | 105.2876428 | 282.0508820 | 44.0850975 | 21.0962017 |
| Dodge Challenger | 163.1150750 | 163.1134210 | 217.7795805 | 72.4403915 | 48.9838851 | 103.4310693 | 103.9023864 | 192.8617917 | 185.8331870 |
| AMC Javelin | 149.6047203 | 149.6014522 | 204.3188913 | 61.3601899 | 61.4274240 | 91.0444349 | 110.3084921 | 180.5479760 | 172.5312555 |
| Camaro Z28 | 233.2228758 | 233.2248748 | 286.0049209 | 163.6632641 | 70.9665308 | 187.8463771 | 10.0761203 | 273.8367985 | 257.7469734 |
| Pontiac Firebird | 248.6780270 | 248.6762035 | 303.3583889 | 156.2240346 | 40.0052475 | 188.5272116 | 80.8057339 | 277.4606884 | 271.3871978 |
| Fiat X1-9 | 92.5048389 | 92.4940020 | 39.8815148 | 184.4471198 | 301.5669483 | 151.4379425 | 333.4843231 | 67.9163981 | 68.5564864 |
| Porsche 914-2 | 44.4033659 | 44.4073589 | 13.1357109 | 139.1579524 | 254.1452553 | 106.0585767 | 285.1986201 | 39.4469276 | 22.1180967 |
| Lotus Europa | 65.7328377 | 65.7362635 | 25.0948550 | 163.2367437 | 272.3582423 | 130.8248192 | 296.4572287 | 72.8971106 | 50.1094030 |
| Ford Pantera L | 245.4247064 | 245.4293785 | 297.2940489 | 180.1140339 | 89.5934049 | 203.0177926 | 21.2655990 | 287.5238795 | 269.9772035 |
| Ferrari Dino | 66.7661029 | 66.7764167 | 90.2415509 | 130.5523007 | 215.0673853 | 106.5694802 | 226.2036333 | 113.3023005 | 80.6550953 |
| Maserati Bora | 265.6454248 | 265.6491465 | 309.7718171 | 229.3419352 | 170.7094473 | 242.4393015 | 107.7224977 | 313.8633093 | 288.8755628 |
| volvo 142E | 39.1894029 | 39.1626037 | 20.6939436 | 137.0363299 | 248.0063378 | 104.1863681 | 275.1353516 | 53.6823481 | 24.6913548 |
| Merc 280 | | | | | | | | | |
| Merc 280C | | | | | | | | | |
| Merc 450SE | | | | | | | | | |
| Merc 450SL | | | | | | | | | |
| Merc 450SLC | | | | | | | | | |
| Cadillac Fleetwood | 315.3904859 | 315.3557081 | 197.8842803 | 197.9154476 | 197.8526242 | | | | |
| Lincoln Continental | 306.6760719 | 306.6406187 | 187.5997191 | 187.6330806 | 187.5671081 | 15.6224446 | | | |
| Chrysler Imperial | 292.7146896 | 292.6989332 | 171.6600758 | 171.6743028 | 171.6537637 | 40.8399636 | 25.3714237 | | |
| Fiat 128 | 106.5053149 | 106.6829794 | 228.3247948 | 228.2592340 | 228.4051825 | 417.7687579 | 410.0206984 | 397.2276375 | |
| Honda Civic | 116.7280991 | 116.8711475 | 238.0141824 | 237.9588183 | 238.0828999 | 425.3271621 | 417.9679574 | 405.8152201 | |
| Toyota Corolla | 113.6290721 | 113.8118009 | 235.5183809 | 235.4481971 | 235.6024098 | 425.3446517 | 417.5429986 | 404.6335386 | |
| Toyota Corona | 54.3641713 | 54.4258314 | 176.6020527 | 176.5727477 | 176.6305359 | 368.3195488 | 360.0267515 | 346.5724649 | |
| Dodge challenger | 152.8929263 | 152.8722437 | 51.8008639 | 51.8242520 | 51.8012606 | 163.6314881 | 156.2805020 | 145.9194779 | |

```
Console Terminal Background Jobs
R 4.2.1 ~\
Merc 280C
Merc 450SE
Merc 450SL
Merc 450SLC
Cadillac Fleetwood
Lincoln continental
Chrysler Imperial
Fiat 128
Honda Civic
Toyota Corolla
Toyota corona
Dodge Challenger
AMC Javelin
Camaro Z28
Pontiac Firebird
Fiat X1-9
Porsche 914-2
Lotus Europa      33.7678653
Ford Pantera L    288.5852993 297.5376920
Ferrari Dino      87.9105966 80.4553451 224.4587490
Maserati Bora     303.9222549 303.2796468 86.9383253 223.5342175
Volvo 142E        18.7555858 27.8104457 277.4803312 70.4751034 289.1157363
> set.seed(240)
> hierarchial_cluster_model <- hclust(mtcars_temp, method = "average")
> hierarchial_cluster_model

call:
hclust(d = mtcars_temp, method = "average")

Cluster method : average
Distance       : euclidean
Number of objects: 32

> # plotting dendrogram
> plot(hierarchial_cluster_model)
> # Cutting tree by height
> abline(h = 110, col = "green")
> # cutting by no. of clusters
> final_fit <- cutree(hierarchial_cluster_model, k = 3 )
> final_fit
      Mazda RX4      Mazda RX4 wag      Datsun 710      Hornet 4 Drive      Hornet Sportabout      Valiant      Duster 360
      1              1              1              2              2              2              2
Merc 240D           Merc 230           Merc 280           Merc 280C           Merc 450SE           Merc 450SL           Merc 450SLC
      1              1              1              1              2              2              2
Cadillac Fleetwood Lincoln Continental Chrysler Imperial      Fiat 128      Honda Civic      Toyota Corolla      Toyota Corona
      2              2              2              1              1              1              1
Dodge Challenger    AMC Javelin      Camaro Z28      Pontiac Firebird      Fiat X1-9      Porsche 914-2      Lotus Europa
      2              2              2              2              1              1              1
Ford Pantera L      Ferrari Dino      Maserati Bora      Volvo 142E
      2              1              3              1

> table(final_fit)
final_fit
 1  2  3
16 15  1
> rect.hclust(hierarchial_cluster_model, k = 3, border = "yellow")
> |
```

